


## Multitasking

# Multitasking

- **Contents**
  - Family life
  - Creating a process from Python
  - Old interface examples
  - Using the subprocess module
  - The subprocess.Popen class
  - Running a basic process
  - Capturing the output
  - Very basic threads in Python
  - Using the multiprocessing module
  - Queue objects
- **Summary**



# Family life

---

- **A process is an instance of a program**
  - Loaded and ready to run
- **Every process has a parent**
  - So every process is a child
- **Child usually inherits attributes of parent**
  - Environment
  - Current directory
  - Open files
- **Relationship depends on the operating system**
  - **UNIX has strong family ties**
    - if a parent dies, the child is an 'orphan'
  - **Microsoft Windows has few ties between parent and child**
    - parent has to explicitly maintain a HANDLE to the child
    - can disown children, but still supports inheritance



Process groups are  
single parent families

# Creating a process from Python

- **Process interfaces can be platform specific**
  - **os.fork**
    - UNIX specific
    - Creates another process - does not run another program
      - **exec***type* is required to run another program after fork
    - Requires os.wait or os.waitpid to avoid a zombie
  - **os.system**
    - Passes the command to the shell (cmd.exe or the Bourne shell)
    - Runs an additional shell process
  - **os.spawn***type*
    - Some types are not available on Windows
    - Can run in a similar mode to **exec***type*
  - **os.popen**
    - Run a process connected through pipes



All these interfaces  
are deprecated

# Old interface examples

- **Run a process and wait for it to complete**
  - Invokes a surrogate shell

```
import os
status = os.system('hello.py')
print "Child exited with", status
```

- **Run a process at the other end of a pipe**
  - Returns a file object

```
for line in os.popen('tasklist').readlines():
    print ":", line,
```



All these interfaces  
are deprecated

# Waiting for a child

---

- **`os.wait()` - UNIX only**
  - No arguments
  - Waits for the child after a fork or spawn
  - Returns a tuple of the child's PID and its exit status
  - Avoids creating a zombie
- **`os.waitpid (PID, options)`**
  - Waits for the child after a fork or spawn
    - *PID*                      Process ID to wait on, -1 to wait on any child
    - *options*                0, or system specific flag, ignored on Windows
  - Return value as `wait()`
  - Avoids creating a zombie

# Using the subprocess module

---

- **Unifying process creation**
  - Introduced at Python 2.4
  - Intended to replace `os.system` and `os.spawn`
- **Main method is Popen**
  - Returns a subprocess object
  - Parameters are discussed over...
- **Other shortcuts are available**
  - `call` and `check_call`
  - `getoutput` and `getstatusoutput` (UNIX specific)
- **We discuss the multiprocessing package later**
  - Runs processes in a similar way to threads

# The subprocess.Popen class

- **Constructor parameters:**

<code>args</code>	Command-line to execute
<code>bufsize=0</code>	Buffersize, 0: unbuffered
<code>executable=None</code>	Program to be executed, rarely needed
<code>stdin=None</code>	Handle to be used for stdin (can be PIPE)
<code>stdout=None</code>	Handle to be used for stdout (can be PIPE)
<code>stderr=None</code>	Handle to be used for stderr (can be STDOUT)
<code>preexec_fn=None</code>	Code to call before the program (UNIX)
<code>close_fds=False</code>	Do not inherit open file handles
<code>shell=False</code>	Use a shell to execute the command
<code>cwd=None</code>	Working directory of the child process
<code>env=None</code>	Environment block of the child process
<code>universal_newlines=False</code>	See any of '\r' or '\n' as newlines
<code>startupinfo=None</code>	Windows only STARTUPINFO struct
<code>creationflags=0</code>	Windows only creation flags



# Running a basic process

- Run a process and wait for it to complete
- A shell is sometimes required
  - When using shell meta-characters
    - Wildcards, pipes, redirections, etc.
  - On Windows, no file association is done unless `shell=True`

```
from subprocess import *  
proc = Popen('hello.py', shell=True)  
proc.wait()  
print "Child exited with",proc.returncode
```

examples  
that follow  
assume this

- Don't use a shell if you don't need to
  - It can add an unnecessary overhead

Typically:  
C:\Python26\python.exe

```
proc = Popen([sys.executable, 'hello.py'])  
proc.wait()
```

# Capturing the output

---

- Use the **communicate** method
  - Returns a tuple of byte objects containing stdout, and stderr

```
proc = Popen('tasklist', stdout=PIPE, stderr=PIPE)
(output, error) = proc.communicate()

if error != None:
    print "error:", error

print "output:", output
```

- Remember that data has to be stored in memory - too much may crash your program!

# Passing data through a pipe

---

- **stdout and stdin are file objects**
- **Read program output one record at a time**
  - Means there is less data held in memory


```
cmd = 'gzip -dc compressed_file.gz'  
pipe = Popen(cmd, stdout=PIPE).stdout  
  
for line in pipe:  
    print ":", line,
```

- **Write to program's input stream one record at a time**

```
cmd = 'lp -'  
proc = Popen(cmd, stdin=PIPE)  
proc.communicate(some_data)
```

# Processes and threads

- **A process is an instance of a running program**
  - It owns a collection of system resources
- **A thread is an asynchronous unit of execution within a process**
  - A process may have several threads
- **Great for spreading processing among processor cores**



# Very basic threads in Python

- Python threading usually uses the **threading** module
  - Call `threading.Thread` (*function*)

```
from threading import Thread
import time

def myfunc(*args):
    print "From thread", args
    time.sleep(5)

th1 = Thread(target=myfunc, args='T1')
th2 = Thread(target=myfunc, args='T2')
th1.start()
th2.start()
print "From main"
th1.join()
th2.join()
```

?

```
From thread ('T'From threadFrom main,
'1'())
T', '2')
```

- Or create our own class derived from `threading.Thread`

# Synchronisation objects in threading

---

- **Several objects are available for thread synchronisation**
  - **Condition variables**
    - Similar to those used by pthreads
  - **Events**
    - Similar to those used by Win32
  - **Thread local storage**
    - Enables global variables to be local to a thread
  - **Locks**
    - Similar to a mutex, has a concept of ownership
  - **Semaphores**
    - A counting lock, e.g. allow 3 threads to access a resource, but no more
  - **Timers**
    - Similar to waitable timers on Win32 and interval timers on UNIX

# Simple use of Lock

- To fix the print issue, and to protect a global list

```
from threading import Lock
csScreen      = Lock()
csSharePrices = Lock()
dSharePrices = []

def GetStockPrice():
    global dSharePrices

    csSharePrices.acquire()
    dPrices = dSharePrices[:]
    csSharePrices.release()
    return dPrices

def Sessions:
    csScreen.acquire()
    print "\nwaiting for requests\n"
    csScreen.release()
```

# The trouble with threads

---

- **They are very difficult to code**
  - Sharing variables requires locking mechanisms
  - Subtle timing differences can make debugging difficult
- **The Python Global Interpreter Lock (GIL)**
  - **The GIL locks the interpreter**
    - Threads are locked for about 100 byte-code instructions
    - Simplifies and protects the interpreter
  - **The GIL *does not* mean that:**
    - Python is not multi-threaded - C modules can multi-thread
    - You don't need to worry about locking - you certainly do!

"Multi-threading is a way of shooting yourself in both feet"



# Using the multiprocessing module

- **Uses processes rather than threads**
  - Default number of processes is one for each core
  - Also supports process pools, and processes across systems
  - Pipes and Queues for synchronised communication

```
from multiprocessing import Process

def myfunc(*args):
    print "From proc", args
    time.sleep(5)

if __name__ == '__main__':
    p1 = Process(target=myfunc, args='T1')
    p2 = Process(target=myfunc, args='T2')
    p1.start()
    p2.start()
    print "From main"
    p1.join()
    p2.join()
```

```
From main
From proc ('T2',)
From proc ('T1',)
```

# Queue objects

- **Used by threads and multiprocessing**
  - Provides a serialised method of communication
  - multiprocessing also supports **JoinableQueue**

```
from multiprocessing import Process, Queue
import os
```

```
def myfunc(*args):
    queue = args[0]

    word = ''
    while word != 'END':
        word = queue.get()
        if len(word) == 7:
            print os.getpid(),":",word
```


Get an item  
from the queue

Continued on next slide...


## Queue objects example(2)

```
if __name__ == '__main__':  
    queue = Queue()  
    p1 = Process(target=myfunc, args=(queue, '1'))  
    p2 = Process(target=myfunc, args=(queue, '2'))  
  
    p1.start()  
    p2.start()  
  
    for line in open('words'):  
        queue.put(line[:-1])  
  
    queue.put ('END')  
    queue.put ('END')  
  
    p1.join()  
    p2.join()  
    print "All done"
```

Put an item  
onto the queue



Make sure there is  
an 'END' marker for  
each child process



# Summary

---

- **Running a program using the older interfaces was platform specific**
  - These functions are now considered deprecated
- **The subprocess module, and the Popen method, provides a more unified approach**
  - Although there are still platform specific methods
- **The communicate method can be used to pass data through pipes**
- **Threads can create more problems than they solve**
- **For true multiprocessing, consider another way**